

Cloud Computing Technologies


DAT515 - Fall 2025

Kubernetes

Prof. Hein Meling



Motivation for Kubernetes



AUTOMATED

KUBERNETES

AUTOMATED
SCALING

CRASHING

SIMPLIFIED
MANAGEMENT

AUTOMATED
SCALING

AUTOMATED
CRASHING

SIMPLIFIED
MANAGEMENT

SELF-HEALING

Challenges with Traditional Mutable Systems

- Incremental Updates
 - Mutable systems apply changes gradually
 - Leads to unpredictable states due to accumulated modifications over time
- Operator Modifications
 - Multiple admin users can modify the system
 - Resulting in undocumented changes and configuration drift

Challenges with Traditional Mutable Systems

- Upgrade Complexity
 - Tools like `apt-get` overwrite old binaries with new ones
 - Directly edit files in `/etc/` without version control
 - Difficult to track changes or roll back effectively
- Unclear System State
 - Current state is not represented as a single, unified artifact
 - Mix of updates and changes, increasing risk of errors and inconsistencies

Tidbits from QuickFeed's Virtual Machine Config

```
sudo setcap CAP_NET_BIND_SERVICE=+eip /path/to/binary/quickfeed
```

```
sudo sysctl net.ipv4.ip_forward=1
```

```
sudo sysctl -p
```

```
sudo service docker restart
```

```
./backup.sh
```

```
./update.sh
```



```
#!/bin/bash

if sysctl -q -n net.ipv4.ip_forward | grep -q '0'; then
    echo "IP packet forwarding disabled on host machine; run these commands"
    echo ""
    echo "sudo sysctl net.ipv4.ip_forward=1"
    echo "sudo sysctl -p"
    echo "sudo service docker restart"
    exit 1
fi

QUICKFEED=$HOME/quickfeed
GOBIN=$HOME/go/bin
DATABASE=qf.db
LOGFILE=quickfeed.log
BACKUP=$QUICKFEED/backups

cd $QUICKFEED

echo "Backing up executables"
cp $GOBIN/quickfeed $BACKUP/quickfeed.$(date +"%m-%d-%y").bak

echo "Backing up the database file"
cp $QUICKFEED/$DATABASE $BACKUP/$DATABASE.$(date +"%m-%d-%y").bak

echo "Backing up the $LOGFILE file"
cp $HOME/logs/$LOGFILE $BACKUP/$LOGFILE.$(date +"%m-%d-%y").bak

echo "Running webpack"
if ! make ui-update; then
    echo "Failed to compile the client"
    exit 1
fi

echo "Running go install"
if ! make install; then
    echo "Failed to compile the server"
    exit 1
fi

if pgrep quickfeed &> /dev/null; then
    echo "Server running; shutting down server..."
    if ! killall quickfeed; then
        echo "Failed to stop the server; trying to restart..."
    fi
fi

echo "Starting the QuickFeed server"
quickfeed &> $HOME/logs/$LOGFILE &

if ! pgrep quickfeed &> /dev/null; then
    echo "Failed to start the server"
    exit 1
fi

echo "All done. QuickFeed server restarted and running"
exit 0
```


Once an artifact is created in the system, it does not change

- Promotes **consistency** and **reliability**
- Immutable Systems
 - New, complete images replace old ones in a single operation
 - **Avoiding incremental changes**
- Example in Containers
 - Instead of updating a running container
 - Build **new container** image enabling **rollback** capability

Once an artifact is created in the system, it does not change

- Reduced Risk of Configuration Drift
 - Immutable containers prevent unrecorded modifications
 - Transparency to changes; easier to troubleshoot
- Antipattern Warning: Imperative **changes** to running containers are **discouraged**, except in extreme, **mission-critical cases**



Mutable

Immutable

Pros

- Does not need to build servers from scratch every time a change is required

- Declarative
- Predictable state
- Auditable
- Easy to rollback
- Modular

Cons

- Incremental changes can fail
- Indiscrete versioning

- Unable to modify in-place
- Misaligned with traditional IT practices

Challenges with Imperative Configuration

- Sequential Execution
 - Imperative commands define actions (e.g., "run A, then B"),
 - Must be executed in order, increasing complexity
- Difficult to Understand Current System State
 - Making debugging harder

Challenges with Imperative Configuration

- Imperative commands are **error-prone**
 - Changes not easily reversible
 - E.g., how do you reverse the steps performed by a shell script
 - Impact may be unclear until after execution
- Imperative instructions more difficult to review in source control
 - Limiting collaboration and rollback options

Specify the Desired State of the System

- State-Based Definition
 - Declarative configurations define the desired state of the system
 - E.g., `replicas = 3`, is clear and easy to understand without execution
 - Less room for error
- Declarative configuration integrated with version control (GitOps)
 - Single source of truth, code review, and testing

Specify the Desired State of the System

- **Easy Rollback**
 - Restating previous declarative state to roll back
 - No need to provide reverse instructions
- Kubernetes **automatically matches** the actual state to the desired state



Continuously take action to retain desired state

- Continuous monitoring and recovery — **guard against failures**
- Reduced Operator Burden
 - Automates repairs, eliminating the need for manual intervention
- Developer Productivity
 - Less maintenance, more time developing new features
- Impact: **Improved system reliability**

What is Kubernetes?



kubernetes

Open-Source Container Orchestration Platform


- Automatically distributes traffic (**load balance**) and exposes containers (**service discovery**)
- Automated (staged) **rollouts** and **rollbacks** (on failures)
- Perform **health checks** and **restart** failed containers
- Automatically **scale** applications up or down based on demand
- Automatically **mounts storage systems** (local or cloud-based)

Open-Source Container Orchestration Platform




kubernetes


- “Kubernetes” is Greek, meaning **helmsman** or **pilot**
- Often abbreviated **k8s**
 - Eight letters between the **k** and the **s**
- Open sourced by Google
- Based on many years of experience with Google’s proprietary Borg systems
 - (Omega)




Once Configured,
Kubernetes Operates Autonomously,
Requiring Minimal Supervision

Recap about Containers







Kubernetes Architecture




- Docker ensures the correct number of containers are running



Kubernetes Cluster Architecture





Kubernetes Cluster Architecture




- **Control plane**
 - Manages worker nodes and Pods in the cluster
 - Production: **Replicated** for fault-tolerance and high availability
- Set of **worker nodes** (machines)
 - Run containerized applications
 - Nodes host **Pods**

Kubernetes Components




Pod — the smallest deployable unit

- Kubernetes manages Pods (not containers)
- **Pod** — group of one or more (tightly coupled) containers
 - **Shared** storage and network resources
 - **Co-located** and **co-scheduled**
 - On same physical or virtual machine
- To run pods, nodes must install a **container runtime**



Pod — the smallest deployable unit

- One-container-per-Pod model
 - Most common
 - Wrapper around a single container
- Multiple containers working together
 - **Shared** storage and network resources
 - **Co-located** and **co-scheduled**



Joke time

With “Pod-casts” of
How do Kubernetes containers communicate?
course!

Because it couldn't handle
Why did the pod break up with the container?
the traffic anymore!

Control Plane Components

Make global decisions about the cluster

- Scheduling
- Detecting and responding to cluster events
 - E.g., start new pod when Deployment's `replicas` field unsatisfied
- Control plane runs on machines separate from worker nodes (typically)


The Kubernetes API server

- The Hub in a Kubernetes cluster
- Exposes the Kubernetes API to
 - Worker nodes
 - Control plane components
 - External systems and tools
- Scale horizontally (deploying more apiserver instances)
 - Load balance traffic between instances



The Kubernetes API server


- Services REST operations to access the cluster's shared state
- Authentication, authorization, and request validation
- Validates and configures data for api objects
 - Pods, services, scheduler, replication controllers, etc



Features

Simple interface

Read and write values using standard HTTP tools, such as curl



Key-value storage


Store data in hierarchically organized directories, as in a standard filesystem

```

/config
├── /database
/feature-flags
├── /verbose-logging
└── /redesign
  
```

Watch for changes

Watch specific keys or directories for changes and react to changes in values



Strongly consistent, distributed key-value store

- Reliable (replicated) data store
- Tolerates **network partitions** and **machine failures**
 - Manual intervention to remain fault-tolerant
- Uses a leader to drive a consensus protocol
 - Raft (see DAT520 Distributed Systems)



Side Track: My PhD Work




Figure 1. Overview of ARM components and interfaces.

```

<Service name="ARM/ReplicationManager">
  <Param name="ServiceMonitorExpiration" value="3"/>
  <ProtocolModules>
    <Module name="GroupMembership"/>
    <Module name="StateMerge"/>
    <Module name="EGMI"/>
    <Module name="Supervision">
      <Param name="GroupFailureSupport" value="no"/>
      <Param name="RemoveDelay" value="5"/>
    </Module>
  </ProtocolModules>
  <DistributionPolicy class="DisperseOnSites"/>
  <ReplicationPolicy class="KeepMinimalInPartition">
    <Redundancy initial="3" minimal="2"/>
  </ReplicationPolicy>
</Service>


```

Figure 8. A sample service configuration description for the RM.

kube-scheduler

Optimize Pod Placement for Efficient Resource Use and Availability

- Monitors unassigned Pods
- Matches Pods to available worker nodes
- Based on resource requirements, constraints, and policies



Optimize Pod Placement for Efficient Resource Use and Availability

- Ensures that workloads are balanced across the cluster
- Scheduling based on factors like
 - resource utilization
 - data locality
 - inter-workload interference
 - affinity/anti-affinity rules
 - hardware/software/policy constraints
 - deadlines and node health

kube-controller-manager

The Controller Concept — from Robotics and Automation

- Control loop — non-terminating loop to regulate the state of a system
 - System has *current state*
 - Specify the *desired state*
- Example: Thermostat
 - Current temperature
 - Specify desired temperature



Controllers Track Kubernetes Resource Type (Object)

- Objects specify desired state
- Controllers are responsible for
 - moving closer to the desired state
- Action: send messages to API server to trigger changes (side effects) to reach desired state



Example Controllers

- Node Controller
 - Detect and respond when nodes go down
- Job Controller
 - Watches for Job objects (one-off tasks)
 - Create Pods to run tasks to completion

A Job is a Kubernetes resource that runs one or more Pods to carry out a task.

cloud-controller-manager

Cloud-specific control logic (optional)

- Runs controllers specific to cloud provider
 - Works similar to kube-controller-manager
- Not used for on-premises or own laptop for learning/testing
- Node Controller
- Route Controller: sets up routes in underlying cloud infrastructure
- Service Controller: creating, updating, and deleting cloud load balancers

Joke time


Because they wanted to

Why did the DevOps engineer bring a ladder to
reach new heights in
the data center?

horizontal scaling!

Node Components

Kubelet



Kubelet

Node Agent — runs on each node

- Register node with apiserver
- Node Resource Monitoring
 - Collects resource usage data like CPU and memory
 - Reports to control plane to help with scheduling and resource management

Node Agent — runs on each node

- Pod Management
 - Monitor pods assigned to kubelet's node
 - **Ensures containers** described in PodSpecs are **running** and **healthy**
- Pod Specification Compliance
 - Specify which containers should be running
 - The kubelet receives pod specifications from control plane

Node Agent — runs on each node

- Container Lifecycle Management
 - Perform container health checks and liveness or readiness probes
 - Start, stop and manage containers
 - Container in a pod crashes — the kubelet can restart it

kube-proxy


Maintains network rules for the node (optional)

- Allow network communication to Pods
 - inside or outside cluster
- Uses OS packet filtering layer (if available)
- Don't need kube-proxy if a network plugin implements packet forwarding

Container Runtime Interface


Enable kubelet to use different container runtimes

- Each node needs a container runtime
- Kubernetes CRI defines
 - gRPC protocol for communication between kubelet and container runtime
 - The kubelet uses CRI to launch Pods and their containers



Alternatives to Docker as Kubernetes runtime

- Kubernetes v0.2 (2014): Docker only supported runtime
- Kubernetes v1.5 (2016): Introduced Container Runtime Interface
- Kubernetes v1.24 (2022): Dropped *built-in* support for Docker



kubectl

Create and manage Kubernetes objects

Management technique	Operates on	Recommended environment	Supported writers	Learning curve
Imperative commands	Live objects	Development projects	1+	Lowest
Imperative object configuration	Individual files	Production projects	1	Moderate
Declarative object configuration	Directories of files	Production projects	1+	Highest



Imperative commands

- `kubectl create deployment nginx --image nginx`
- Advantages:
 - Single word action
 - Single step to make changes to cluster
- Disadvantages:
 - No change review process and audit trail

Imperative object configuration

- `kubectl create -f nginx.yaml`
- Advantages:
 - Config stored in git
 - Change review process and audit trail
- Disadvantages:
 - Must understand object schema (yaml format)
 - Updates must be reflected in config files, or they will be lost

Declarative object configuration

- `kubectl diff -f configs/`
`kubectl apply -f configs/`
- Advantages:
 - Support operating on directories
 - Automatically detect operation type (create, patch, delete) per object
- Disadvantages:
 - Harder to debug and understand
 - Partial updates mean lead to complex merge and patch operations

Joke time

Because they always lose

Why do Kubernetes nodes make terrible
their connection to the
comedians?

audience!

Scaling Your Service and Your Teams

Decouple Components from each other

- Component Isolation
 - Make components function independent of each other
 - Define API for component to interact with other (consumer) components
 - Hide internal details
- Allows scaling of individual services without impacting other components


Team Efficiency

- APIs enable dev teams to focus on smaller, manageable microservices
- Reducing the need for cross-team coordination
- Clear API boundaries lowers the communication demands across teams
- Faster development and deployment

Monolithic vs Microservices Architecture




Monolithic Architecture



Microservices Architecture





Questions?

- <https://www.cncf.io/blog/2019/08/19/how-kubernetes-works/>
-